

Keep It Simple: Testing Databases via Differential Query Plans


Jinsheng Ba, Manuel Rigger
National University of Singapore


Background: Logic Bugs

Logic bugs refer to incorrect results returned by DBMSs.

t0	t1
c0	c0
2	0


```
SELECT COUNT(*)
FROM t0 INNER JOIN
t1 ON t0.c0==t1.c0;
```

0 

1 

Study

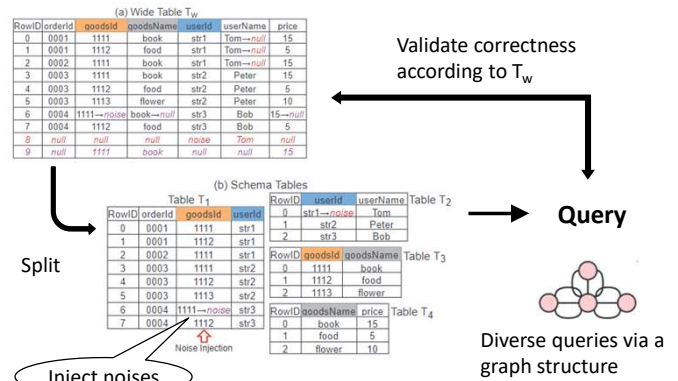
We found that most bugs found by TQS can be reproduced by comparing the execution results of the same query of different query plans.

```
SELECT t0.c0 FROM t0 WHERE t0.c0 IN
(SELECT t0.c0 FROM t0 WHERE
(t0.c0 NOT IN
(SELECT t0.c0 FROM t0 WHERE t0.c0 )
) = (t0.c0)
); -- {0000001985} , {0000001996}
SELECT t0.c0 FROM t0 WHERE t0.c0 IN
(SELECT /*+ no_semi_join() */ t0.c0 FROM t0 WHERE
(t0.c0 NOT IN
(SELECT t0.c0 FROM t0 WHERE t0.c0 )
) = (t0.c0)
); -- empty set 
```

Problem and Challenges

How to automatically find logic bugs?

A state-of-the-art work: Transformed Query Synthesis (TQS)



Challenges:

1. TQS is a **sophisticated** method that requires splitting tables and retrieving results from the first table.
2. This method can **only be applied** to equijoin.

Method: Differential Query Plans (DQP)

Comparing query plan executions is a simple testing method, and we demonstrated that such a simple method can be as effective as TQS. We propose the method **Differential Query Plans (DQP)**:

① Database State Generation

user	transaction
user_id	transaction_id amount
1	1_c12934 100000
2	1_e3b664 -10



② Query Generation

```
SELECT
IFNULL(SUM(amount), 0) AS balance
FROM user JOIN transaction
ON transaction.transaction_id =
user.user_id;
```


③ Query Plan Enforcement

```
SELECT
/*+ JOIN_ORDER(transaction, user) */
IFNULL(SUM(amount), 0) as balance
FROM user JOIN transaction
ON transaction.transaction_id =
user.user_id;
```

Query Plan
nested_loop
+ table
table_name: user
access_type: index
+ table
table_name: transaction
access_type: all

④ Result Validation


balance
99990.00



balance
0.00

Results

In MySQL, MariaDB, and TiDB, DQP can reproduce 14 of 15 bugs found by TQS. Additionally, DQP found 26 new bugs.



Conclusion

1. DQP is a **general black-box** method, which is easy to understand.
2. DQP is **lightweight** as it was implemented in less than 100 lines of Java code.
3. DQP is **applicable** as at least 8 of 10 most popular relational DBMSs support enforcing query plans.
4. DQP is a **simple alternative** to TQS that achieves the same level of effectiveness.



Paper



Code

Session 35: Security (3)
Thursday June 13
5:15 pm – 5:30 pm
Location: Puyehue/Calbuco



Author: Jinsheng Ba



Author: Manuel Rigger